

Deduce, Propose, Correct: Probability-Accountable LLM Guidance for Typed Program Inference

Tri Nguyen and Dat Nguyen

Abstract

Large language models can guide program synthesis, but their practical value is difficult to separate from symbolic pruning, test execution, and selection. Moreover, sequential Monte Carlo requires the probability of each proposal, whereas free-form code generation exposes no tractable probability for a canonical program. We introduce DPC, a controlled framework for typed programming by example in a bounded functional language. Type-directed generalization proposes program skeletons; sound deduction refutes inconsistent skeletons and derives examples for typed holes; and a language model scores finite canonical choices. A defensive mixture of language-model energies, deduction, and uniform support yields an explicit proposal that includes family, hole, construction-trace, and cloning probabilities and can therefore be importance-corrected. Exact enumeration supplies an oracle for finite-target diagnostics, while a separate lazy path measures online program evaluations. Exploratory experiments reveal an important negative result: raw full-prompt Qwen energies often prefer short incorrect abstract syntax trees, while deduction rescues hard transformations. The framework is designed to make such failures measurable rather than to treat any passing sample as evidence that an LLM improved synthesis.

Keywords: program synthesis, programming by example, large language models, sequential Monte Carlo, neuro-symbolic inference

1 Introduction

Programming by example (PBE) asks for a program that maps each supplied input to its paired output. The specification is accessible, but it is incomplete: many programs may fit a small example set, and the space of typed recursive programs grows combinatorially. Classical systems address this tension with domain languages, structural hypotheses, deduction, and an Occam-style search order (Gulwani, 2011; Feser et al., 2015). Recent systems instead use learned policies or large language models (LLMs) to prioritize plausible code (Kalyan et al., 2018; Barke et al., 2024; Li et al., 2024).

The empirical question “did the LLM help?” is nevertheless easy to answer incorrectly. A system can attribute success to an LLM even when a type checker, deduction rule, exhaustive catalog construction, or post-hoc selector already identified the solution. An LLM can also assign high probability to a compact but semantically wrong serialization. Reporting only the best passing sample hides both effects. A useful scientific instrument must isolate the LLM’s proposal from the symbolic support, target score, and selection mechanism.

Sequential Monte Carlo (SMC) provides a natural population view of synthesis: programs are particles, semantic fit defines a potential, and resampling moves computation toward promising regions. This view appears in neural program synthesis (Ellis et al., 2019) and, more recently, in ModelSMC’s inference over executable scientific models (Wahl

et al., 2026). Its probabilistic interpretation, however, depends on the proposal probability $q(e' \mid e)$. For free-form LLM output, many token strings may denote the same abstract syntax tree (AST), invalid generations induce fallback mass, and common APIs do not return the probability of the canonical AST actually consumed by the search.

We study a narrower setting in which this proposal-density gap can be closed. The program language, constants, structural cost, and hole catalogs are bounded. The LLM does not generate arbitrary code; it assigns teacher-forced energies to a complete finite set of canonical family or hole choices. The application normalizes these energies, combines them with a sound deduction guide and a uniform floor, and records the probability of every selected construction step. This produces an evaluable proposal without pretending that it is the free-form probability of an AST.

Our contributions are: (i) a type-directed factorized construction space that combines structural generalization, sound refutation, and derived hole examples; (ii) a canonical finite LLM-energy proposal with explicit family, sequential-hole, construction-trace, and clone accounting; (iii) an importance-corrected Feynman–Kac target whose materialized finite controls can be checked against exact enumeration; and (iv) an evaluation protocol that separates uniform, deduction-only, LLM-only, and combined proposals and preserves negative results. We do not claim that the current bounded language is a general synthesizer, that finite- N particles are a calibrated posterior, or that full enumeration is faster than search.

2 Background and Problem Setting

Let $\mathcal{D} = \{(x_j, y_j)\}_{j=1}^M$ be examples with an inferred signature $\tau_{\text{in}} \rightarrow \tau_{\text{out}}$. A program e is exact on $\mathcal{D}$ when $e(x_j) = y_j$ for every j . We work with a typed functional DSL containing integer and Boolean expressions, lists, conditionals, `map`, and right folds. A structural cost $C(e)$ prefers compact combinator programs, following the minimum-cost perspective of Lambda² (Feser et al., 2015). Exactness on $\mathcal{D}$ is not a proof of the unknown intended semantics, so evaluation also uses held-out inputs generated before confirmatory runs.

For retained finite construction-trace support $\mathcal{E}_{\mathcal{D}}$, we define a data-conditional base measure. A trace $z = (h, c_{1:K})$ records a family and one choice for each typed hole, and assembles to program $e(z)$. If $h(z) \in \mathcal{H}$ denotes its family, then

$$p_0(z \mid \mathcal{D}) = \frac{1}{|\mathcal{H}|} \frac{\exp[-\lambda_C C(e(z))]}{\sum_{u \in \mathcal{E}_{h(z)}} \exp[-\lambda_C C(e(u))]} \quad (1)$$

The equal family mass prevents a large catalog from dominating solely through cardinality. Because deduction and bounded-catalog policy depend on $\mathcal{D}$, p_0 is a conditional Occam measure, not an unconditional Bayesian prior.

At inverse temperature β , the finite Gibbs target is

$$\tilde{\pi}_{\beta}(z) = \mathbf{1}[z \in \mathcal{E}_{\mathcal{D}}] p_0(z \mid \mathcal{D}) \exp[-\beta \lambda_L L_{\mathcal{D}}(e(z))], \quad (2)$$

where $L_{\mathcal{D}}$ is a deterministic soft execution loss. This is a declared finite target for controlled inference. It is not a posterior over all programs and the loss exponential is not asserted to be a real-world likelihood.

3 Typed Generalization and Deduction

Generalization maps the inferred signature and structural relations in the examples to typed hypotheses. For list-to-list tasks these include a direct expression, a map, a generic fold, and factorized filter-map folds. Each hypothesis is a program skeleton with typed holes. For example,

$$\text{foldr}(\lambda i, a. \text{ if } p(i) \text{ then } g(i) :: a \text{ else } a, [], x) \quad (3)$$

has holes $p : \text{Int} \rightarrow \text{Bool}$ and $g : \text{Int} \rightarrow \text{Int}$.

Deduction has two roles. Refutation removes a skeleton only when an example contradicts an invariant; for instance, `map` is impossible when input and output lengths differ. Example inference transfers top-level evidence to a hole: singleton and aligned list examples produce facts such as $p(-3) = \text{false}$ or $g(2) = 4$. Ambiguous alignments do not generate facts. Deduction therefore changes proposal guidance while preserving all candidates not ruled out by a sound invariant.

Each surviving hole receives a complete bounded catalog of well-typed canonical ASTs. Generic holes are enumerated by structural cost; declared factorized skeletons use explicit finite catalogs. The online state is the family-plus-hole construction trace, so two traces that assemble to the same AST remain distinct latent states. The materialized canonical-program control may instead apply an explicit owner rule and records the discarded-alias count. We compare its oracle metrics with online traces only when the construction map is injective on the reported support; otherwise the two modes define different targets.

4 Canonical LLM and Deduction Proposals

At a construction node A with finite choices $\mathcal{C}_A$, Qwen receives a common prompt containing the task, typed hole, sound deductions, current ancestor, and previous fillings. For each canonical choice c , vLLM returns teacher-forced prompt-token log probabilities for the complete concatenated text. These values define an application-level finite energy; they are not interpreted as a continuation probability or the probability of a free-form AST. We archive the prompt prefix, candidate strings, token identifiers, per-token scores, energy semantics, and declared model and tokenizer revision labels so that the Qwen categorical can be re-played. Specifically, if $\ell_\theta(A, c)$ is the sum of the scored conditional token log probabilities for the complete concatenated prompt and $n(A, c)$ is its number of scored positions, the implementation supports $s_\theta(A, c) = \ell_\theta(A, c)$ and $\bar{s}_\theta(A, c) = \ell_\theta(A, c)/n(A, c)$. Both are full-prompt energies; the mean is not a candidate-suffix likelihood because no separate tokenizer boundary between the textual prefix and candidate is assumed.

Let $s_\theta(A, c)$ be the declared Qwen energy and let $r_A(c)$ be the exact subtree marginal of an Occam-weighted deduction guide. The categorical used for both sampling and correction is

$$q_A(c) = (1 - \varepsilon) \left[(1 - \rho) \text{softmax} \left(\frac{s_\theta(A, \cdot)}{\tau} \right)_c + \rho r_A(c) \right] + \frac{\varepsilon}{|\mathcal{C}_A|}. \quad (4)$$

The uniform term gives every declared choice positive mass. Mixing normalized distributions, rather than adding token scores to logical penalties, prevents long JSON serializations from overwhelming deduction merely through cumulative log probability.

A program trace z selects one family and an ordered sequence of hole fillings:

$$q_{\text{construct}}(z \mid a, \mathcal{D}) = q_H(h \mid a, \mathcal{D}) \prod_k q_{A_k}(c_k \mid a, h, c_{<k}, \mathcal{D}). \quad (5)$$

When only one family survives, $q_H = 1$ and no provider call is made. Identical requests created by resampling are scored once and fanned out without changing the distribution.

With clone probability $\alpha < 1$, the transition on construction traces is

$$Q_\alpha(z' \mid z) = \alpha \mathbf{1}[z' = z] + (1 - \alpha) q_{\text{construct}}(z' \mid z, \mathcal{D}). \quad (6)$$

For $z' = z$, both routes are combined with a log-sum-exp. Provider errors abort the accountable mode because an unmodeled ancestor fallback would add unknown mass. A run-wide score budget is reserved before each request wave.

5 Importance-Corrected SMC

For stages $t = 1, \dots, T$, define the path target $\gamma_t(z_{1:t}) = \prod_{s=1}^t \tilde{\pi}_{\beta_s}(z_s)$ and transition $M_t = Q_\alpha$. The incremental potential is

$$G_t(z_{t-1}, z_t) = \frac{\tilde{\pi}_{\beta_t}(z_t)}{Q_\alpha(z_t \mid z_{t-1}, \mathcal{D})}. \quad (7)$$

Thus $M_t G_t = \tilde{\pi}_{\beta_t}$ and the normalized terminal-state marginal is the declared finite target. If resampling occurs, ancestor weights reset to $1/N$; otherwise the previous normalized weights remain the base.

The product of stage normalizers belongs to the path target, not merely the final Gibbs target. Exact enumeration on small supports supplies the reference $\sum_t \log Z_{\beta_t}$ and the final target distribution. We report total-variation distance, exact-program mass, mean loss and cost, effective sample size (ESS), and normalizer error. These diagnostics measure a finite particle approximation; they do not imply that a small particle set is itself calibrated.

The implementation exposes two execution boundaries. The reference mode materializes and evaluates every complete bounded program so that target functionals are known exactly. The online mode first samples initial traces and then proposal traces, assembling and evaluating only previously unvisited samples. It reports search cost but cannot report exact TV or normalizer error unless a matching external oracle is run separately. Oracle comparisons require the same trace semantics; any online-oracle comparison must first verify that cross-family aliases do not change the state space. This separation prevents exhaustive support construction from being counted as an LLM or SMC synthesis success.

6 Experimental Protocol

Our primary factorial comparison has four proposal arms: uniform (U), deduction plus uniform (D), Qwen (Q), and Qwen plus deduction (Q+D). All arms share type inference, sound refutation, finite catalogs, target, seeds, particles, iterations, and proposal-score caps. Paired contrasts D–U, Q–U, Q+D–D, and Q+D–Q isolate the soft guides. The current protocol does not include a separate uncorrected arm or a one-stage versus multi-stage resampling ablation; either would require an additional pre-registered matrix.

Table 1: Controlled enumerable tasks; counts are fixed regression-test targets.

Task	Examples	Support	Exact states	Max choices/path
Integer sign predicate	5	82	1	82
Map increment	4	16,012	3	392
Fold sum	5	22,914	46	316
Positive filter	4	24,646	144	518
Bounded square	14	36,198	2	663
Signed window	14	132,198	4	823

Each run archives exact training consistency, held-out semantic correctness, per-stage newly and cumulatively realized traces, total finite candidate choices scored, wall time, ESS, resampling, final particles, and failures. The current aggregator exposes training and held-out outcomes, loss, cost, ESS, reference diagnostics when available, total candidate scores, and wall time; it does not yet derive weight variance or executions and scores to first success. The seed, rather than a particle, is the independent unit. The draft protocol defines paired intention-to-treat exact discovery as the primary endpoint, orders the arm contrasts for multiplicity reporting, and counts failed provider runs as failures. It does not yet freeze or implement confidence intervals, hypothesis tests, or adjusted p -values; these and any missing secondary endpoint must be implemented before the protocol is frozen and before a result is called confirmatory.

Table 1 summarizes the current controlled benchmark. The two no-solution support diagnostics are excluded from success rates and reported separately. These tasks exercise scalar expressions, map, fold, filter-map, and signed piecewise mapping. They remain a small synthetic DSL benchmark; external PBE or SyGuS-derived tasks are required before making a broad synthesis claim.

7 Exploratory Pilot Findings

Table 2 reports matched exploratory, oracle-backed runs, not confirmatory estimates or online speed measurements. The hardest signed-window support contains 132,198 states and four exact ASTs. With four particles and one stage, the deduction-guided arms assigned terminal particle mass to an exact program, while uniform and Qwen-only did not. On bounded square, Qwen-only selected no exact program whereas the combined arm selected two. These outcomes measure proposal behavior after support construction and motivated the pre-registered repeated-seed protocol; they are not used to estimate population success rates or faster synthesis.

The score decomposition supplies a sharper negative result. In the exact signed-window particle, Qwen assigned probabilities approximately 0.0034 to the correct family, 9.3×10^{-10} to the selected exact window predicate, and 1.0×10^{-11} to the exact piecewise mapper. The deduction component assigned approximately 1.0, 0.392, and 0.419, respectively. Therefore the combined exact selection is evidence for the defensive symbolic mixture, not evidence that Qwen discovered the hard program.

Table 2: Exploratory oracle-backed seed-23 pilots. “Scores” counts finite candidate choices scored; TV compares the four-particle terminal measure with exact enumeration. Neither column measures online synthesis speed.

Task	Proposal	Exact	Best loss	Scores	TV
Signed window	Uniform	no	20	1,688	1.000
Signed window	Deduction + uniform	yes	0	2,490	0.543
Signed window	Qwen	no	20	886	1.000
Signed window	Qwen + deduction	yes	0	1,688	0.847
Bounded square	Qwen	no	29	128	1.000
Bounded square	Qwen + deduction	yes	0	2,652	0.636
Map increment	Qwen	no	6	212	1.000
Map increment	Qwen + deduction	no	6	838	1.000

Raw full-prompt sequence sums strongly favor short AST serializations. On map increment, Qwen-only placed nearly all weighted mass on the empty-list expression, and the four-particle combined run also missed the exact map. We therefore pre-register energy-semantic ablations rather than silently replacing the score after inspecting results. Total full-prompt energy is the historical baseline, and mean full-prompt energy is the implemented full-path length-normalization ablation. Explicitly token-bounded continuation energy remains future work; it requires a provider/tokenizer protocol that establishes the boundary before it can be evaluated as a declared finite energy with its own exact q .

8 Related Work

Lambda² combines type-aware inductive generalization, deductive refutation, derived examples, and best-first enumeration for recursive data structures (Feser et al., 2015). Neural-guided deductive search and DreamCoder learn distributions that prioritize symbolic search (Kalyan et al., 2018; Ellis et al., 2021). HYSYNTH derives a context-free surrogate from LLM completions, and LLM-guided enumeration iteratively exchanges search state with a model (Barke et al., 2024; Li et al., 2024). Our narrower focus is the exact probability accounting of a canonical factorized energy proposal and its separation from symbolic deduction and target scoring.

Write–Execute–Assess combines learned program policies with SMC (Ellis et al., 2019); ModelSMC represents executable models as particles and uses LLM refinement (Wahl et al., 2026). SMC also steers token generation under probabilistic-program constraints (Lew et al., 2023; Zhao et al., 2024). We do not claim the first use of SMC or LLMs in synthesis. The contribution is a finite setting in which family and hole proposal terms, defensive mixtures, and clone routes are all evaluable and can be checked against an exact target.

9 Limitations and Broader Impact

The DSL and current tasks are deliberately small. Specialized filter-map catalogs encode useful structural knowledge, support depends on observed data, and multiple training-exact programs may disagree on unseen inputs. Exact enumeration scales poorly and is an oracle, not the online algorithm. The semantic core is extensively tested but not formally verified. Model and tokenizer revisions can change Qwen energies, and full-prompt scoring has a serialization-length confound. These boundaries rule out claims of general program synthesis, formal correctness, or a posterior over unbounded source code.

The lazy implementation avoids executing complete programs that are never sampled, but current pilots do not establish lower wall time or fewer total resources because symbolic catalog construction and Qwen scoring also have cost. The method can also amplify biases encoded in prompts or the model. A uniform floor preserves mathematical support, not social fairness or semantic safety. Generated programs must remain sandboxed, resource-bounded, and reviewed before deployment. We publish failed runs and complete score ledgers so that readers can audit adverse model preferences rather than only successful outputs.

10 Conclusion

LLM-guided synthesis needs more than a plausible generated program: it needs a clear account of what the model proposed, what symbolic reasoning removed or favored, and how selection changed the resulting measure. DPC turns typed PBE into a finite construction problem whose LLM energies, deduction guide, uniform support, and importance correction are separately inspectable. The current pilots show why that decomposition matters: Qwen can harm search, and deduction can rescue it. Confirmatory benchmarks must determine where the LLM adds value and where classical search remains preferable.

Acknowledgments and Disclosure of Funding

This draft received no declared external funding. Compute was run on a rented RunPod NVIDIA A100 instance. The final submission must replace this statement with complete funding, donated-compute, competing-interest, and author disclosures.

Appendix A. Proposal Normalization

For every node A , Equation 4 is a convex combination of three normalized categorical distributions and therefore sums to one. Because $\varepsilon > 0$, each choice has mass at least $\varepsilon/|C_A|$. Products over the finite construction trie are normalized by induction on trie depth. Equation 6 is a mixture of a point mass and this normalized construction distribution, so it is also normalized and has full support over the declared finite trace space.

Appendix B. Claim Boundary for Exact Enumeration

Reference enumeration discovers whether exact programs exist while computing the oracle. Consequently, an oracle-backed run demonstrates probability accounting and estimator diagnostics, not faster synthesis. The current pilot table is explicitly oracle-backed and makes no online discovery claim. Any future online result must use the lazy path, count every newly realized trace from initialization and proposal stages, and remain separate from the oracle. An exact state first observed during reference construction is never credited to Qwen or SMC discovery.

References

- Shraddha Barke, Emmanuel Anaya Gonzalez, Saketh Ram Kasibatla, Taylor Berg-Kirkpatrick, and Nadia Polikarpova. HYSYNTH: Context-free LLM approximation for guiding program synthesis. In *Advances in Neural Information Processing Systems*, volume 37, 2024. URL <https://openreview.net/forum?id=5jt0ZSA6Co>.
- Kevin Ellis, Maxwell Nye, Yewen Pu, Felix Sosa, Joshua Tenenbaum, and Armando Solar-Lezama. Write, execute, assess: Program synthesis with a REPL. In *Advances in Neural Information Processing Systems*, volume 32, 2019. URL <https://arxiv.org/abs/1906.04604>.
- Kevin Ellis, Catherine Wong, Maxwell Nye, Mathias Sablé-Meyer, Lucas Morales, Luke Hewitt, Luc Cary, Armando Solar-Lezama, and Joshua B. Tenenbaum. DreamCoder: Bootstrapping inductive program synthesis with wake-sleep library learning. In *Proceedings of the 42nd ACM SIGPLAN International Conference on Programming Language Design and Implementation*, pages 835–850. ACM, 2021. doi: 10.1145/3453483.3454080.
- John K. Feser, Swarat Chaudhuri, and Isil Dillig. Synthesizing data structure transformations from input-output examples. In *Proceedings of the 36th ACM SIGPLAN Conference on Programming Language Design and Implementation*, pages 229–239. ACM, 2015. doi: 10.1145/2737924.2737977.
- Sumit Gulwani. Automating string processing in spreadsheets using input-output examples. In *Proceedings of the 38th ACM SIGPLAN-SIGACT Symposium on Principles of Programming Languages*, pages 317–330. ACM, 2011. doi: 10.1145/1926385.1926423.
- Ashwin Kalyan, Abhishek Mohta, Oleksandr Polozov, Dhruv Batra, Prateek Jain, and Sumit Gulwani. Neural-guided deductive search for real-time program synthesis from examples. In *International Conference on Learning Representations*, 2018. URL <https://openreview.net/forum?id=rywDjg-RW>.
- Alexander K. Lew, Zhi-Xuan Tan, Gabriel Grand, and Vikash K. Mansinghka. Sequential monte carlo steering of large language models using probabilistic programs. *arXiv preprint arXiv:2306.03081*, 2023. doi: 10.48550/arXiv.2306.03081.
- Yixuan Li, Julian Parsert, and Elizabeth Polgreen. Guiding enumerative program synthesis with large language models. *arXiv preprint arXiv:2403.03997*, 2024. doi: 10.48550/arXiv.2403.03997.

Stefan Wahl, Raphaela Schenk, Ali Farnoud, Jakob H. Macke, and Daniel Gedon. A probabilistic framework for LLM-based model discovery. In *Proceedings of the 43rd International Conference on Machine Learning*, volume 306 of *Proceedings of Machine Learning Research*, 2026. URL <https://arxiv.org/abs/2602.18266>.

Stephen Zhao, Rob Brekelmans, Alireza Makhzani, and Roger Grosse. Probabilistic inference in language models via twisted sequential monte carlo. *arXiv preprint arXiv:2404.17546*, 2024. doi: 10.48550/arXiv.2404.17546.